

Visual Reasoning

Mohan Bhambhani

Topics in Deep Learning, 2017

Outline

1 Introduction

2 Datasets

3 Models

Introduction

Visual reasoning vs VQA

- Too much bias in VQA datasets.

Introduction

Visual reasoning vs VQA

- Too much bias in VQA datasets.
- Maybe **What covers the ground?** is answerable not because it understands the scene but because biased datasets often ask questions about the ground when it is **snow-covered**.

Introduction

Visual reasoning vs VQA

- Too much bias in VQA datasets.
- Maybe **What covers the ground?** is answerable not because it understands the scene but because biased datasets often ask questions about the ground when it is **snow-covered**.
- **Solution:** Focus on the ultimate goal of visual reasoning.

Introduction


Visual reasoning vs VQA

- Too much bias in VQA datasets.
- Maybe **What covers the ground?** is answerable not because it understands the scene but because biased datasets often ask questions about the ground when it is **snow-covered**.
- **Solution:** Focus on the ultimate goal of visual reasoning.
- **Goal of Visual systems:** To completely understand information contained in the images.

Introduction

Visual reasoning vs VQA

- Too much bias in VQA datasets.
- Maybe **What covers the ground?** is answerable not because it understands the scene but because biased datasets often ask questions about the ground when it is **snow-covered**.
- **Solution:** Focus on the ultimate goal of visual reasoning.
- **Goal of Visual systems:** To completely understand information contained in the images.



Outline

1 Introduction


2 Datasets

3 Models


- To avoid bias added by humans, questions must be generated automatically. But, at the same time they must be good.

- To avoid bias added by humans, questions must be generated automatically. But, at the same time they must be good.
- Getting this done is easier in synthetic images.

- To avoid bias added by humans, questions must be generated automatically. But, at the same time they must be good.
- Getting this done is easier in synthetic images.
- Interface:




Sample chain-structured question:




What color is the cube to the right of the yellow sphere?

Sample tree-structured question:



How many cylinders are in front of the small thing and on the left side of the green object?

CLEVR function catalog



Images:

- Synthetically generated

Images:

- Synthetically generated
- 3 object shapes (cube, sphere, and cylinder)

Images:

- Synthetically generated
- 3 object shapes (cube, sphere, and cylinder)
- 2 sizes (small and large)

Images:

- Synthetically generated
- 3 object shapes (cube, sphere, and cylinder)
- 2 sizes (small and large)
- 2 materials (shiny 'metal' and matte 'rubber')

Images:

- Synthetically generated
- 3 object shapes (cube, sphere, and cylinder)
- 2 sizes (small and large)
- 2 materials (shiny 'metal' and matte 'rubber')
- 8 colors

Images:

- Synthetically generated
- 3 object shapes (cube, sphere, and cylinder)
- 2 sizes (small and large)
- 2 materials (shiny 'metal' and matte 'rubber')
- 8 colors
- 4 relationships between objects: ('left', 'right', 'behind', and 'in front')

Images:

- Synthetically generated
- 3 object shapes (cube, sphere, and cylinder)
- 2 sizes (small and large)
- 2 materials (shiny 'metal' and matte 'rubber')
- 8 colors
- 4 relationships between objects: ('left', 'right', 'behind', and 'in front')
- Scene graph annotated

Images:

- Synthetically generated
- 3 object shapes (cube, sphere, and cylinder)
- 2 sizes (small and large)
- 2 materials (shiny 'metal' and matte 'rubber')
- 8 colors
- 4 relationships between objects: ('left', 'right', 'behind', and 'in front')
- Scene graph annotated
- **To generate image:** Randomly sample a scene graph. Then generate image using Blender (a tool to generate 3d images from scene graph).

Questions:

Questions:

- Define question families.

Questions:

- Define question families.
- Example: How many red things are there? can be formed from How many $\langle C \rangle$ $\langle M \rangle$ things are there?, binding the parameters $\langle C \rangle$ and $\langle M \rangle$ (with types color and material) with values red and nil.

Questions:

- Define **question families**.
- Example: **How many red things are there?** can be formed from **How many $\langle C \rangle \langle M \rangle$ things are there?**, binding the parameters $\langle C \rangle$ and $\langle M \rangle$ (with types **color** and **material**) with values **red** and **nil**.
- Each question in CLEVR is associated with a functional program that can be executed on an image's scene graph.

Questions:

- Define **question families**.
- Example: **How many red things are there?** can be formed from **How many $\langle C \rangle \langle M \rangle$ things are there?**, binding the parameters $\langle C \rangle$ and $\langle M \rangle$ (with types **color** and **material**) with values **red** and **nil**.
- Each question in CLEVR is associated with a functional program that can be executed on an image's scene graph.
- The functional program `count(filterColor(red, scene()))` for this question can be formed by instantiating the associated program template `count(filterColor($\langle C \rangle$, filterMaterial($\langle M \rangle$, scene())))`.

Questions:

- Define **question families**.
- Example: **How many red things are there?** can be formed from **How many $\langle C \rangle \langle M \rangle$ things are there?**, binding the parameters $\langle C \rangle$ and $\langle M \rangle$ (with types **color** and **material**) with values **red** and **nil**.
- Each question in CLEVR is associated with a functional program that can be executed on an image's scene graph.
- The functional program `count(filterColor(red, scene()))` for this question can be formed by instantiating the associated program template `count(filterColor($\langle C \rangle$, filterMaterial($\langle M \rangle$, scene())))`.
- CLEVR contains a total of 90 question families, each with a single program template and an average of four text templates.

Questions:

- Define **question families**.
- Example: **How many red things are there?** can be formed from **How many $\langle C \rangle \langle M \rangle$ things are there?**, binding the parameters $\langle C \rangle$ and $\langle M \rangle$ (with types **color** and **material**) with values **red** and **nil**.
- Each question in CLEVR is associated with a functional program that can be executed on an image's scene graph.
- The functional program `count(filterColor(red, scene()))` for this question can be formed by instantiating the associated program template `count(filterColor($\langle C \rangle$, filterMaterial($\langle M \rangle$, scene())))`.
- CLEVR contains a total of 90 question families, each with a single program template and an average of four text templates.
- So to generate question: choose a question family, select values for each of its template parameters, execute the resulting program on the image's scene graph to find the answer.

Quality of questions :

- Still many ill-posed questions can be formed.

Quality of questions :

- Still many ill-posed questions can be formed.
- The question **What color is the cube to the right of the sphere?** would be ill-posed if there were many cubes right of the sphere, or degenerate if there were only one cube in the scene.


Quality of questions :

- Still many ill-posed questions can be formed.
- The question **What color is the cube to the right of the sphere?** would be ill-posed if there were many cubes right of the sphere, or degenerate if there were only one cube in the scene.
- Employ a depth-first search. At each step of assigning values to the variables, use ground-truth scene information to prune nodes which are guaranteed to produce undesirable questions.

Quality of questions :

- Still many ill-posed questions can be formed.
- The question **What color is the cube to the right of the sphere?** would be ill-posed if there were many cubes right of the sphere, or degenerate if there were only one cube in the scene.
- Employ a depth-first search. At each step of assigning values to the variables, use ground-truth scene information to prune nodes which are guaranteed to produce undesirable questions.
- Maintain uniform question distribution across families.

CLEVR




```

"params": [
  {"type": "Size", "name": "<Z>"},
  {"type": "Color", "name": "<C>"},
  {"type": "Material", "name": "<M>"},
  {"type": "Shape", "name": "<S>"},
  {"type": "Relation", "name": "<R>"},
  {"type": "Size", "name": "<Z2>"},
  {"type": "Color", "name": "<C2>"},
  {"type": "Material", "name": "<M2>"},
  {"type": "Shape", "name": "<S2>"}
],
"text": [
  "What size is the <Z2> <C2> <M2> <S2> [that is] <R> the <Z> <C> <M> <S>?",
  "What is the size of the <Z2> <C2> <M2> <S2> [that is] <R> the <Z> <C> <M> <S>?",
  "How big is the <Z2> <C2> <M2> <S2> [that is] <R> the <Z> <C> <M> <S>?",
  "There is a <Z2> <C2> <M2> <S2> [that is] <R> the <Z> <C> <M> <S>; what size is it?",
  "There is a <Z2> <C2> <M2> <S2> [that is] <R> the <Z> <C> <M> <S>; how big is it?",
  "There is a <Z2> <C2> <M2> <S2> [that is] <R> the <Z> <C> <M> <S>; what is its size?"
],

```


Split	Images	Questions	Unique questions	Overlap with train
Total	100,000	999,968	853,554	-
Train	70,000	699,989	608,607	-
Val	15,000	149,991	140,448	17,338
Test	15,000	149,988	140,352	17,335



Outline

1 Introduction

2 Datasets

3 Models

- **Inspiration:** Different architectures are needed for different tasks - CNNs for object detection, RNNs for counting.

- **Inspiration:** Different architectures are needed for different tasks - CNNs for object detection, RNNs for counting.
- Decompose a question into its linguistic substructures and train a neural network module for each substructure.

- **Inspiration:** Different architectures are needed for different tasks - CNNs for object detection, RNNs for counting.
- Decompose a question into its linguistic substructures and train a neural network module for each substructure.
- Jointly train the modules and dynamically compose them into deep networks which can learn to answer the question.

- **Inspiration:** Different architectures are needed for different tasks - CNNs for object detection, RNNs for counting.
- Decompose a question into its linguistic substructures and train a neural network module for each substructure.
- Jointly train the modules and dynamically compose them into deep networks which can learn to answer the question.
- Start by analyzing the question and decide what logical units are needed to answer the question and what should be the relationship between them.

- Find: Finds objects of interest.



- Transform: Shift regions of attention.



- Combine: Merge two attention maps into a single one.



- Describe: Map a pair of attention and input image to a distribution over the labels.



- Measure: Map attention to a distribution over the labels.



- Model is specified by collection of modules $\{m\}$ and a network layout predictor P which maps from strings to networks.

- Model is specified by collection of modules $\{m\}$ and a network layout predictor P which maps from strings to networks.
- Map question to the layout which specifies the set of modules and connections between them.

- Model is specified by collection of modules $\{m\}$ and a network layout predictor P which maps from strings to networks.
- Map question to the layout which specifies the set of modules and connections between them.
- Assemble the final network using the layout.





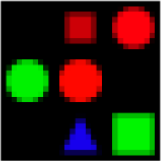
- Model is specified by collection of modules $\{m\}$ and a network layout predictor P which maps from strings to networks.
- Map question to the layout which specifies the set of modules and connections between them.
- Assemble the final network using the layout.
- Parse the input question to obtain set of dependencies and obtain a representation similar to combinatory logic.

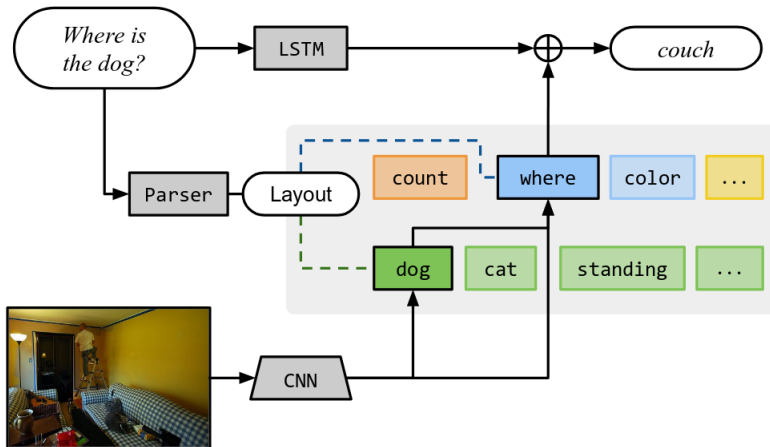
- Model is specified by collection of modules $\{m\}$ and a network layout predictor P which maps from strings to networks.
- Map question to the layout which specifies the set of modules and connections between them.
- Assemble the final network using the layout.
- Parse the input question to obtain set of dependencies and obtain a representation similar to combinatory logic.
- For example, **what is standing in the field** becomes **what(stand)**, **what color is the truck** becomes **color(truck)**, and **is there a circle next to a square** becomes **is(circle, next-to(square))**.

- Model is specified by collection of modules $\{m\}$ and a network layout predictor P which maps from strings to networks.
- Map question to the layout which specifies the set of modules and connections between them.
- Assemble the final network using the layout.
- Parse the input question to obtain set of dependencies and obtain a representation similar to combinatory logic.
- For example, **what is standing in the field** becomes **what(stand)**, **what color is the truck** becomes **color(truck)**, and **is there a circle next to a square** becomes **is(circle, next-to(square))**.
- The symbolic representation is mapped to a layout:
 - All leaves become find module.
 - All internal nodes become transform/combine module.
 - All root nodes become describe/measure module.

- Final model combines the output from the NMN with predictions from a simple LSTM question encoder. This helps in modelling the syntactic and semantic regularities of the question.

- Final model combines the output from the NMN with predictions from a simple LSTM question encoder. This helps in modelling the syntactic and semantic regularities of the question.

				
<i>how many different lights in various different shapes and sizes?</i>	<i>what is the color of the horse?</i>	<i>what color is the vase?</i>	<i>is the bus full of passengers?</i>	<i>is there a red shape above a circle?</i>
<code>describe[count](find[light])</code>	<code>describe[color](find[horse])</code>	<code>describe[color](find[vase])</code>	<code>describe[is](combine[and](find[bus], find[full])</code>	<code>measure[is](combine[and](find[red], transform[above](find[circle])))</code>
four (four)	brown (brown)	green (green)	yes (yes)	yes (yes)



Implementation details:

- Since some modules are updated more frequently than others, adaptive per weight learning rates are better.
- Training: Adadelta with standard parameter settings.
- Stanford parser has F1 score of 87%.

Model	Overall	Count	Exist	Compare numbers	Query	Compare attribute
<i>Human</i>	92.6	86.7	96.6	86.5	95.0	96.0
<i>LSTM</i>	46.8	41.7	61.1	69.8	36.8	51.8
<i>CNN+RNN</i>	52.3	43.7	65.2	67.1	49.3	53.0
<i>SA</i>	68.5	52.2	71.1	73.5	85.3	52.3
<i>NMN</i>	72.1	79.3	52.5	71.4	78.9	78

- NMNs are restricted to parsers for model configuration. A better model would be that learns it from the data.
- This paper learns to both parse the language into linguistic structures and compose them into appropriate layouts.
- Model has 2 parts. First, a set of co-attentive neural modules that provide parameterized functions for solving sub-tasks, and a layout policy to predict a question-specific layout from which a neural network is dynamically assembled.

N2NMNs

- x_{vis} - image feature map, x_{txt} - text features

N2NMNs

- x_{vis} - image feature map, x_{txt} - text features

N2NMs

- x_{vis} - image feature map, x_{txt} - text features
- **Find:** to localize some objects and output an attention map over the image. $a_{out} = conv2(conv1(x_{vis}) \odot Wx_{txt})$

N2NMNs

- x_{vis} - image feature map, x_{txt} - text features
- **Find:** to localize some objects and output an attention map over the image. $a_{out} = conv2(conv1(x_{vis}) \odot Wx_{txt})$
- **Relocate:** to transform the input image attention map and output a new attention map

$$a_{out} = conv2(conv1(x_{vis}) \odot W_1 sum(a \odot x_{vis}) \odot W_2 x_{txt})$$

N2NMNs

- x_{vis} - image feature map, x_{txt} - text features
- **Find:** to localize some objects and output an attention map over the image. $a_{out} = conv2(conv1(x_{vis}) \odot Wx_{txt})$
- **Relocate:** to transform the input image attention map and output a new attention map
 $a_{out} = conv2(conv1(x_{vis}) \odot W_1sum(a \odot x_{vis}) \odot W_2x_{txt})$
- **And:** To take and of 2 feature maps $a_{out} = minimum(a_1, a_2)$

- x_{vis} - image feature map, x_{txt} - text features
- **Find:** to localize some objects and output an attention map over the image. $a_{out} = conv2(conv1(x_{vis}) \odot Wx_{txt})$
- **Relocate:** to transform the input image attention map and output a new attention map
 $a_{out} = conv2(conv1(x_{vis}) \odot W_1 sum(a \odot x_{vis}) \odot W_2 x_{txt})$
- **And:** To take and of 2 feature maps $a_{out} = minimum(a_1, a_2)$
- **Or:** To take or of 2 feature maps $a_{out} = maximum(a_1, a_2)$

- x_{vis} - image feature map, x_{txt} - text features
- **Find:** to localize some objects and output an attention map over the image. $a_{out} = conv2(conv1(x_{vis}) \odot Wx_{txt})$
- **Relocate:** to transform the input image attention map and output a new attention map
 $a_{out} = conv2(conv1(x_{vis}) \odot W_1 sum(a \odot x_{vis}) \odot W_2 x_{txt})$
- **And:** To take and of 2 feature maps $a_{out} = minimum(a_1, a_2)$
- **Or:** To take or of 2 feature maps $a_{out} = maximum(a_1, a_2)$
- **Filter:** Added this to simplify layout. $a_{out} = and(a, find[x_{vis}, x_{txt}]())$

- x_{vis} - image feature map, x_{txt} - text features
- **Find:** to localize some objects and output an attention map over the image. $a_{out} = conv2(conv1(x_{vis}) \odot Wx_{txt})$
- **Relocate:** to transform the input image attention map and output a new attention map
 $a_{out} = conv2(conv1(x_{vis}) \odot W_1 sum(a \odot x_{vis}) \odot W_2 x_{txt})$
- **And:** To take and of 2 feature maps $a_{out} = minimum(a_1, a_2)$
- **Or:** To take or of 2 feature maps $a_{out} = maximum(a_1, a_2)$
- **Filter:** Added this to simplify layout. $a_{out} = and(a, find[x_{vis}, x_{txt}]())$
- **Exist, Count:** For simple inference from attention map.
 $y = W^T vec(a)$

- x_{vis} - image feature map, x_{txt} - text features
- **Find:** to localize some objects and output an attention map over the image. $a_{out} = \text{conv2}(\text{conv1}(x_{vis}) \odot Wx_{txt})$
- **Relocate:** to transform the input image attention map and output a new attention map
 $a_{out} = \text{conv2}(\text{conv1}(x_{vis}) \odot W_1 \text{sum}(a \odot x_{vis}) \odot W_2 x_{txt})$
- **And:** To take and of 2 feature maps $a_{out} = \text{minimum}(a_1, a_2)$
- **Or:** To take or of 2 feature maps $a_{out} = \text{maximum}(a_1, a_2)$
- **Filter:** Added this to simplify layout. $a_{out} = \text{and}(a, \text{find}[x_{vis}, x_{txt}]())$
- **Exist, Count:** For simple inference from attention map.
 $y = W^T \text{vec}(a)$
- **Describe:** For complex inference $y = W_1^T (W_2 \text{sum}(a \odot x_{vis}) \odot W_3 x_{txt})$

- x_{vis} - image feature map, x_{txt} - text features
- **Find:** to localize some objects and output an attention map over the image. $a_{out} = conv2(conv1(x_{vis}) \odot Wx_{txt})$
- **Relocate:** to transform the input image attention map and output a new attention map
 $a_{out} = conv2(conv1(x_{vis}) \odot W_1sum(a \odot x_{vis}) \odot W_2x_{txt})$
- **And:** To take and of 2 feature maps $a_{out} = minimum(a_1, a_2)$
- **Or:** To take or of 2 feature maps $a_{out} = maximum(a_1, a_2)$
- **Filter:** Added this to simplify layout. $a_{out} = and(a, find[x_{vis}, x_{txt}]())$
- **Exist, Count:** For simple inference from attention map.
 $y = W^T vec(a)$
- **Describe:** For complex inference $y = W_1^T (W_2sum(a \odot x_{vis}) \odot W_3x_{txt})$
- **Eq-count, More, Less:** For pairwise comparison over action maps.
 $y = W_1^T vec(a_1) + W_2^T vec(a_2)$

- x_{vis} - image feature map, x_{txt} - text features
- **Find:** to localize some objects and output an attention map over the image. $a_{out} = conv2(conv1(x_{vis}) \odot Wx_{txt})$

- **Relocate:** to transform the input image attention map and output a new attention map

$$a_{out} = conv2(conv1(x_{vis}) \odot W_1 sum(a \odot x_{vis}) \odot W_2 x_{txt})$$

- **And:** To take and of 2 feature maps $a_{out} = minimum(a_1, a_2)$
- **Or:** To take or of 2 feature maps $a_{out} = maximum(a_1, a_2)$
- **Filter:** Added this to simplify layout. $a_{out} = and(a, find[x_{vis}, x_{txt}]())$
- **Exist, Count:** For simple inference from attention map.

$$y = W^T vec(a)$$

- **Describe:** For complex inference $y = W_1^T (W_2 sum(a \odot x_{vis}) \odot W_3 x_{txt})$
- **Eq-count, More, Less:** For pairwise comparison over action maps.

$$y = W_1^T vec(a_1) + W_2^T vec(a_2)$$

- **Compare:** For complex pairwise comparison

$$y = W_1^T (W_2 sum(a_1 \odot x_{vis}) \odot W_3 sum(a_2 \odot x_{vis}) \odot W_4 x_{txt})$$


Layout:

- To predict the layout
represent the layout in term
of Reverse Polish notation
and convert it to a sequence.

layout
expression

`eq_count(find(), and(find(), find()))`

syntax tree



Reverse Polish
Notation

`[find, find, find, and, eq_count]`


Layout:

- To predict the layout
represent the layout in term
of Reverse Polish notation
and convert it to a sequence.
- This can now be addressed
using Encoder-decoder
framework with both
encoder and decoder being
LSTMs.

layout
expression

eq_count(find(), and(find(), find()))

syntax tree



Reverse Polish
Notation

[find, find, find, and, eq_count]


Layout:

- To predict the layout
represent the layout in term
of Reverse Polish notation
and convert it to a sequence.
- This can now be addressed
using Encoder-decoder
framework with both
encoder and decoder being
LSTMs.
- At each time step in the decoder LSTM, a soft attention map over
the input sequence is predicted.

layout
expression

eq_count(find(), and(find(), find()))

syntax tree



Reverse Polish
Notation

[find, find, find, and, eq_count]


Layout:

- To predict the layout
represent the layout in term
of Reverse Polish notation
and convert it to a sequence.
- This can now be addressed
using Encoder-decoder
framework with both
encoder and decoder being
LSTMs.
- At each time step in the decoder LSTM, a soft attention map over
the input sequence is predicted.
- In the end of each cell of decoder we have softmax layer to predict
the module.

layout
expression

eq_count(find(), and(find(), find()))

syntax tree



Reverse Polish
Notation

[find, find, find, and, eq_count]

Layout:


- To predict the layout
represent the layout in term
of Reverse Polish notation
and convert it to a sequence.

- This can now be addressed
using Encoder-decoder
framework with both
encoder and decoder being
LSTMs.

layout
expression

eq_count(find(), and(find(), find()))

syntax tree



Reverse Polish
Notation

[find, find, find, and, eq_count]

- At each time step in the decoder LSTM, a soft attention map over the input sequence is predicted.
- In the end of each cell of decoder we have softmax layer to predict the module.
- At each step the text features passed to the module are the attended features obtained from the LSTM

Training:

- Training is done end to end.

Training:

- Training is done end to end.
- Cross entropy loss function is used for all the modules.

Training:

- Training is done end to end.
- Cross entropy loss function is used for all the modules.
- Loss function is not fully differentiable since the layout is discrete.

Training:

- Training is done end to end.
- Cross entropy loss function is used for all the modules.
- Loss function is not fully differentiable since the layout is discrete.
- Policy gradient is used for layout selection network.

Training:

- Training is done end to end.
- Cross entropy loss function is used for all the modules.
- Loss function is not fully differentiable since the layout is discrete.
- Policy gradient is used for layout selection network.
- q- question, Im- image, l -layout.

Training:

- Training is done end to end.
- Cross entropy loss function is used for all the modules.
- Loss function is not fully differentiable since the layout is discrete.
- Policy gradient is used for layout selection network.
- q- question, Im- image, l -layout.
- Layout network: $p(l|q; \theta)$

Training:

- Training is done end to end.
- Cross entropy loss function is used for all the modules.
- Loss function is not fully differentiable since the layout is discrete.
- Policy gradient is used for layout selection network.
- q- question, Im- image, l -layout.
- Layout network: $p(l|q; \theta)$
- Loss: $L(\theta) = E_{l \sim p(l|q; \theta)}[\tilde{L}(\theta, l; q, Im)]$

Training:

- Training is done end to end.
- Cross entropy loss function is used for all the modules.
- Loss function is not fully differentiable since the layout is discrete.
- Policy gradient is used for layout selection network.
- q- question, Im- image, l -layout.
- Layout network: $p(l|q; \theta)$
- Loss: $L(\theta) = E_{l \sim p(l|q; \theta)}[\tilde{L}(\theta, l; q, Im)]$
- Gradient from PG:
$$\nabla_{\theta} L = E_{l \sim p(l|q; \theta)}[\tilde{L}(\theta, l) \nabla_{\theta} \log p(l|q; \theta) + \nabla_{\theta} \tilde{L}(\theta, l)]$$

Training:

- Training is done end to end.
- Cross entropy loss function is used for all the modules.
- Loss function is not fully differentiable since the layout is discrete.
- Policy gradient is used for layout selection network.
- q- question, Im- image, l -layout.
- Layout network: $p(l|q; \theta)$
- Loss: $L(\theta) = E_{l \sim p(l|q; \theta)} [\tilde{L}(\theta, l; q, Im)]$
- Gradient from PG:

$$\nabla_{\theta} L = E_{l \sim p(l|q; \theta)} [\tilde{L}(\theta, l) \nabla_{\theta} \log p(l|q; \theta) + \nabla_{\theta} \tilde{L}(\theta, l)]$$
- This is estimated using Monte-Carlo sampling. (M=1)

$$\nabla_{\theta} L \sim \frac{1}{M} \sum_{m=1}^M (E_{l \sim p(l|q; \theta)} [\tilde{L}(\theta, l) \nabla_{\theta} \log p(l|q; \theta) + \nabla_{\theta} \tilde{L}(\theta, l)])$$

Training:

- Training is done end to end.
- Cross entropy loss function is used for all the modules.
- Loss function is not fully differentiable since the layout is discrete.
- Policy gradient is used for layout selection network.
- q- question, Im- image, l -layout.
- Layout network: $p(l|q; \theta)$
- Loss: $L(\theta) = E_{l \sim p(l|q; \theta)}[\tilde{L}(\theta, l; q, Im)]$
- Gradient from PG:

$$\nabla_{\theta} L = E_{l \sim p(l|q; \theta)}[\tilde{L}(\theta, l) \nabla_{\theta} \log p(l|q; \theta) + \nabla_{\theta} \tilde{L}(\theta, l)]$$
- This is estimated using Monte-Carlo sampling. (M=1)

$$\nabla_{\theta} L \sim \frac{1}{M} \sum_{m=1}^M (E_{l \sim p(l|q; \theta)}[\tilde{L}(\theta, l) \nabla_{\theta} \log p(l|q; \theta) + \nabla_{\theta} \tilde{L}(\theta, l)])$$
- To reduce variance of estimated gradient baseline is added.

Training:

- Training is done end to end.
- Cross entropy loss function is used for all the modules.
- Loss function is not fully differentiable since the layout is discrete.
- Policy gradient is used for layout selection network.
- q- question, Im- image, l -layout.
- Layout network: $p(l|q; \theta)$
- Loss: $L(\theta) = E_{l \sim p(l|q; \theta)}[\tilde{L}(\theta, l; q, Im)]$
- Gradient from PG:

$$\nabla_{\theta} L = E_{l \sim p(l|q; \theta)}[\tilde{L}(\theta, l) \nabla_{\theta} \log p(l|q; \theta) + \nabla_{\theta} \tilde{L}(\theta, l)]$$
- This is estimated using Monte-Carlo sampling. (M=1)

$$\nabla_{\theta} L \sim \frac{1}{M} \sum_{m=1}^M (E_{l \sim p(l|q; \theta)}[\tilde{L}(\theta, l) \nabla_{\theta} \log p(l|q; \theta) + \nabla_{\theta} \tilde{L}(\theta, l)])$$
- To reduce variance of estimated gradient baseline is added.
- Layout is pretrained with KL-divergence between expert policy and layout policy loss.

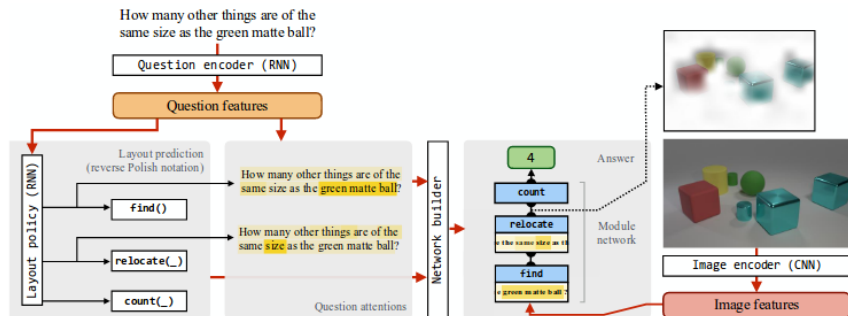
Training:

- Training is done end to end.
- Cross entropy loss function is used for all the modules.
- Loss function is not fully differentiable since the layout is discrete.
- Policy gradient is used for layout selection network.
- q- question, Im- image, l -layout.
- Layout network: $p(l|q; \theta)$
- Loss: $L(\theta) = E_{l \sim p(l|q; \theta)} [\tilde{L}(\theta, l; q, Im)]$
- Gradient from PG:

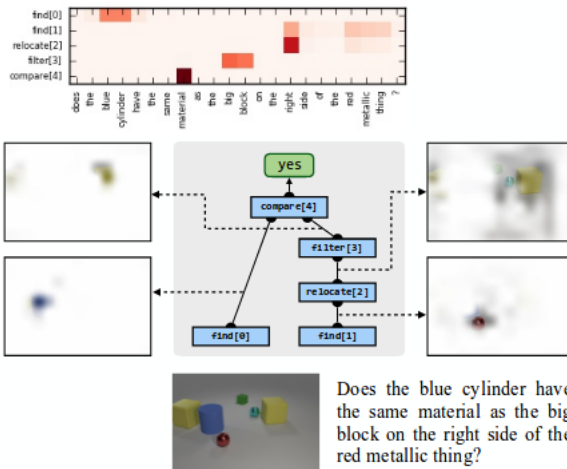
$$\nabla_{\theta} L = E_{l \sim p(l|q; \theta)} [\tilde{L}(\theta, l) \nabla_{\theta} \log p(l|q; \theta) + \nabla_{\theta} \tilde{L}(\theta, l)]$$
- This is estimated using Monte-Carlo sampling. (M=1)

$$\nabla_{\theta} L \sim \frac{1}{M} \sum_{m=1}^M (E_{l \sim p(l|q; \theta)} [\tilde{L}(\theta, l) \nabla_{\theta} \log p(l|q; \theta) + \nabla_{\theta} \tilde{L}(\theta, l)])$$
- To reduce variance of estimated gradient baseline is added.
- Layout is pretrained with KL-divergence between expert policy and layout policy loss.
- Adam optimiser was used for training.

Architecture overview:



Example:



Model	Overall	Count	Exist	Compare numbers	Query	Compare attribute
<i>Human</i>	92.6	86.7	96.6	86.5	95.0	96.0
<i>LSTM</i>	46.8	41.7	61.1	69.8	36.8	51.8
<i>CNN+RNN</i>	52.3	43.7	65.2	67.1	49.3	53.0
<i>SA</i>	68.5	52.2	71.1	73.5	85.3	52.3
<i>NMN</i>	72.1	79.3	52.5	71.4	78.9	78
<i>N2NMN</i>	83.7	85.7	68.5	83.7	90	88.7

- Very similar architecture and same training procedure.

Program generator

- Very similar architecture and same training procedure.
- More modules

Program generator

- Very similar architecture and same training procedure.
- More modules
- Unary: Scene(to get action map), filter-color, filter-shape, filter-material, filter-size, unique, query-color, query-shape, query-material, query-color, query-size, exist, relate, count

Program generator

- Very similar architecture and same training procedure.
- More modules
- Unary: Scene(to get action map), filter-color, filter-shape, filter-material, filter-size, unique, query-color, query-shape, query-material, query-color, query-size, exist, relate, count
- Binary: intersect, union, equal-size, equal-color, equal-material, equal-shape, equal-integer, less-than, greaterthan.

- Very similar architecture and same training procedure.
- More modules
- Unary: Scene(to get action map), filter-color, filter-shape, filter-material, filter-size, unique, query-color, query-shape, query-material, query-color, query-size, exist, relate, count
- Binary: intersect, union, equal-size, equal-color, equal-material, equal-shape, equal-integer, less-than, greaterthan.
- All the modules have residual layers.

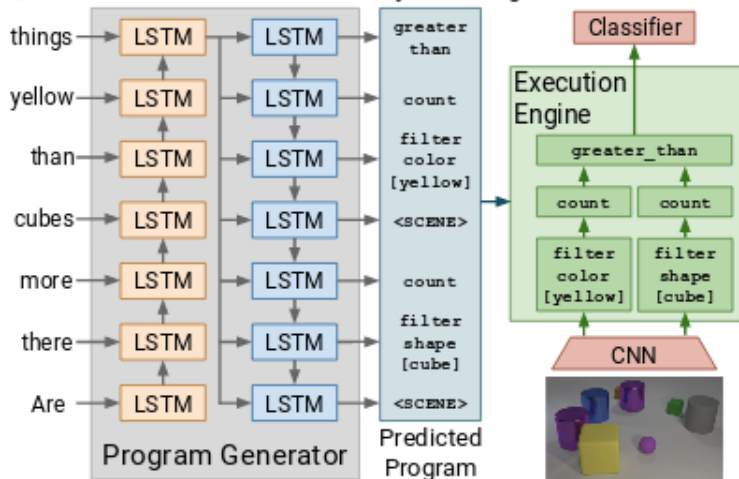
- Very similar architecture and same training procedure.
- More modules
- Unary: Scene(to get action map), filter-color, filter-shape, filter-material, filter-size, unique, query-color, query-shape, query-material, query-color, query-size, exist, relate, count
- Binary: intersect, union, equal-size, equal-color, equal-material, equal-shape, equal-integer, less-than, greaterthan.
- All the modules have residual layers.
- At the end there is output module with fully connected layers.

Program generator


- Very similar architecture and same training procedure.
- More modules
- Unary: Scene(to get action map), filter-color, filter-shape, filter-material, filter-size, unique, query-color, query-shape, query-material, query-color, query-size, exist, relate, count
- Binary: intersect, union, equal-size, equal-color, equal-material, equal-shape, equal-integer, less-than, greaterthan.
- All the modules have residual layers.
- At the end there is output module with fully connected layers.
- Program generation is top-down.

Program generator

Question: Are there more cubes than yellow things? **Answer:** Yes




Program generator



Q: Is there a blue box
in the items? **A:** yes

Predicted Program:
`exist`
`filter_shape[cube]`
`filter_color[blue]`
`scene`


Predicted Answer:
✓ yes



Q: What shape object
is farthest right?
A: cylinder

Predicted Program:
`query_shape`
`unique`
`relate[right]`
`unique`
`filter_shape[cylinder]`
`filter_color[blue]`
`scene`

Predicted Answer:
✓ cylinder



Q: Are all the balls small?
A: no

Predicted Program:
`equal_size`
`query_size`
`unique`
`filter_shape[sphere]`
`scene`
`query_size`
`unique`
`filter_shape[sphere]`
`filter_size[small]`
`scene`
Predicted Answer:
✓ no

Program generator

Model	Overall	Count	Exist	Compare numbers	Query	Compare attribute
<i>Human</i>	92.6	86.7	96.6	86.5	95.0	96.0
<i>LSTM</i>	46.8	41.7	61.1	69.8	36.8	51.8
<i>CNN+RNN</i>	52.3	43.7	65.2	67.1	49.3	53.0
<i>SA</i>	68.5	52.2	71.1	73.5	85.3	52.3
<i>NMN</i>	72.1	79.3	52.5	71.4	78.9	78
<i>N2NMN</i>	83.7	85.7	68.5	83.7	90	88.7
<i>PG</i>	96.9	97.1	92.7	98.6	98.1	98.9

Relational networks

- Symbolic approaches to artificial intelligence are inherently relational.

Relational networks

- Symbolic approaches to artificial intelligence are inherently relational.
- In deep learning, many approaches often struggle in data-poor problems where the underlying structure is characterized by sparse but complex relations.


Relational networks

- Symbolic approaches to artificial intelligence are inherently relational.
- In deep learning, many approaches often struggle in data-poor problems where the underlying structure is characterized by sparse but complex relations.
- This paper proposes relational network for relational reasoning with neural networks.

Relational networks


- Symbolic approaches to artificial intelligence are inherently relational.
- In deep learning, many approaches often struggle in data-poor problems where the underlying structure is characterized by sparse but complex relations.
- This paper proposes relational network for relational reasoning with neural networks.

Original Image:




Non-relational question:

What is the size of the brown sphere?



Relational question:

Are there any rubber things that have the same size as the yellow metallic cylinder?



Relational networks

- An RN is a neural network module with a structure primed for relational reasoning.

Relational networks

- An RN is a neural network module with a structure primed for relational reasoning.
- The design philosophy behind RNs is to constrain the functional form of a neural network so that it captures the core common properties of relational reasoning.

Relational networks

- An RN is a neural network module with a structure primed for relational reasoning.
- The design philosophy behind RNs is to constrain the functional form of a neural network so that it captures the core common properties of relational reasoning.
- $RN(O) = f_{\phi}(\sum_{i,j} g_{\theta}(o_i, o_j))$

Relational networks

- An RN is a neural network module with a structure primed for relational reasoning.
- The design philosophy behind RNs is to constrain the functional form of a neural network so that it captures the core common properties of relational reasoning.
- $RN(O) = f_{\phi}(\sum_{i,j} g_{\theta}(o_i, o_j))$
- The input here is set of objects.

Relational networks

- An RN is a neural network module with a structure primed for relational reasoning.
- The design philosophy behind RNs is to constrain the functional form of a neural network so that it captures the core common properties of relational reasoning.
- $RN(O) = f_{\phi}(\sum_{i,j} g_{\theta}(o_i, o_j))$
- The input here is set of objects.
- f_{ϕ} and g_{θ} are MLPs, and the parameters are learnable synaptic weights, making RNs end-to-end differentiable.

Relational networks

- An RN is a neural network module with a structure primed for relational reasoning.
- The design philosophy behind RNs is to constrain the functional form of a neural network so that it captures the core common properties of relational reasoning.
- $RN(O) = f_{\phi}(\sum_{i,j} g_{\theta}(o_i, o_j))$
- The input here is set of objects.
- f_{ϕ} and g_{θ} are MLPs, and the parameters are learnable synaptic weights, making RNs end-to-end differentiable.
- **RNs learn to infer relations** The functional form in Equation shows that an RN should consider the potential relations between all object pairs.

Relational networks

- An RN is a neural network module with a structure primed for relational reasoning.
- The design philosophy behind RNs is to constrain the functional form of a neural network so that it captures the core common properties of relational reasoning.
- $RN(O) = f_{\phi}(\sum_{i,j} g_{\theta}(o_i, o_j))$
- The input here is set of objects.
- f_{ϕ} and g_{θ} are MLPs, and the parameters are learnable synaptic weights, making RNs end-to-end differentiable.
- **RNs learn to infer relations** The functional form in Equation shows that an RN should consider the potential relations between all object pairs.
- Not knowing the actual relations, RNs must learn to infer the existence of relations.

- **RNs are data efficient** RNs use a single function g_θ to compute each relation.

- **RNs are data efficient** RNs use a single function g_θ to compute each relation.
- This mode of operation encourages greater generalization for computing relations, since g_θ is encouraged not to over-fit to the features of any particular object pair.

- **RNs are data efficient** RNs use a single function g_θ to compute each relation.
- This mode of operation encourages greater generalization for computing relations, since g_θ is encouraged not to over-fit to the features of any particular object pair.
- There will be $O(n^2)$ passes over g_θ . At the same times each example has so many forward passes to learn.

- **RNs are data efficient** RNs use a single function g_θ to compute each relation.
- This mode of operation encourages greater generalization for computing relations, since g_θ is encouraged not to over-fit to the features of any particular object pair.
- There will be $O(n^2)$ passes over g_θ . At the same times each example has so many forward passes to learn.
- **RNs operate on a set of objects** RNs are order invariant of object sequence. This invariance ensures that the RN's output contains information that is generally representative of the relations that exist in the object set.

Relational networks

- In CLEVR, many questions are explicitly relational in nature. Remarkably, powerful QA architectures are unable to solve CLEVR, presumably because they cannot handle core relational aspects of the task.

Model:

Relational networks

- In CLEVR, many questions are explicitly relational in nature. Remarkably, powerful QA architectures are unable to solve CLEVR, presumably because they cannot handle core relational aspects of the task.

Model:

- Get $d \times d$ feature map from CNN.

Relational networks

- In CLEVR, many questions are explicitly relational in nature. Remarkably, powerful QA architectures are unable to solve CLEVR, presumably because they cannot handle core relational aspects of the task.

Model:

- Get $d \times d$ feature map from CNN.
- Each of the d^2 k -dimensional cells in the $d \times d$ feature maps was tagged with an arbitrary coordinate indicating its relative spatial position, and was treated as an object for the RN.

- In CLEVR, many questions are explicitly relational in nature. Remarkably, powerful QA architectures are unable to solve CLEVR, presumably because they cannot handle core relational aspects of the task.

Model:

- Get $d \times d$ feature map from CNN.
- Each of the d^2 k -dimensional cells in the $d \times d$ feature maps was tagged with an arbitrary coordinate indicating its relative spatial position, and was treated as an object for the RN.
- The existence and meaning of an object-object relation should be question dependent.

Relational networks

- In CLEVR, many questions are explicitly relational in nature. Remarkably, powerful QA architectures are unable to solve CLEVR, presumably because they cannot handle core relational aspects of the task.

Model:

- Get $d \times d$ feature map from CNN.
- Each of the d^2 k -dimensional cells in the $d \times d$ feature maps was tagged with an arbitrary coordinate indicating its relative spatial position, and was treated as an object for the RN.
- The existence and meaning of an object-object relation should be question dependent.
- So the RN architecture is modified such that g_θ could condition its processing on the question: $a = f_\phi(\sum_{i,j} g_\theta(o_i, o_j, q))$.

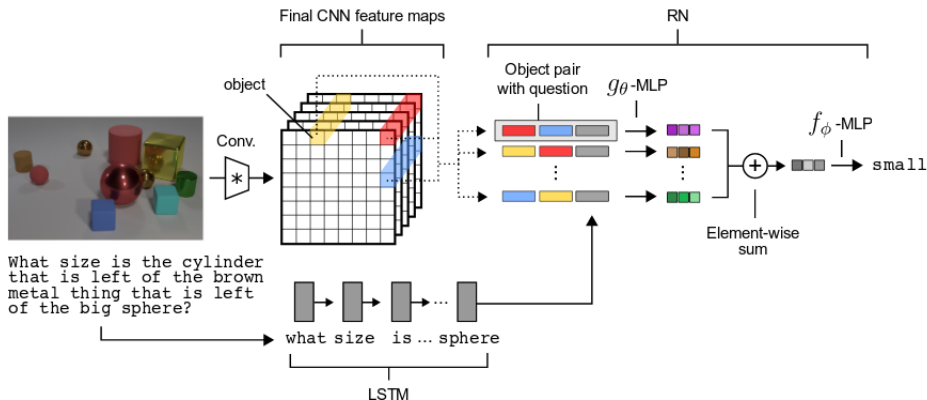
Relational networks

- In CLEVR, many questions are explicitly relational in nature. Remarkably, powerful QA architectures are unable to solve CLEVR, presumably because they cannot handle core relational aspects of the task.

Model:

- Get $d \times d$ feature map from CNN.
- Each of the d^2 k -dimensional cells in the $d \times d$ feature maps was tagged with an arbitrary coordinate indicating its relative spatial position, and was treated as an object for the RN.
- The existence and meaning of an object-object relation should be question dependent.
- So the RN architecture is modified such that g_θ could condition its processing on the question: $a = f_\phi(\sum_{i,j} g_\theta(o_i, o_j, q))$.
- LSTM is used to get question embedding.

Relational networks



- CNN- 4 convolutional layers each with 24 kernel

Relational networks

- CNN- 4 convolutional layers each with 24 kernel
- 128 unit LSTM for question processing, 32 unit word-lookup embedding

Relational networks

- CNN- 4 convolutional layers each with 24 kernel
- 128 unit LSTM for question processing, 32 unit word-lookup embedding
- 4-layer MLP consisting of 256 units per layer with ReLU non-linearities for g_{θ} .

Relational networks

- CNN- 4 convolutional layers each with 24 kernel
- 128 unit LSTM for question processing, 32 unit word-lookup embedding
- 4-layer MLP consisting of 256 units per layer with ReLU non-linearities for g_θ .
- 3-layer MLP consisting of 256, 256 (with 50% dropout), and 29 units with ReLU non-linearities for f_ϕ .

- CNN- 4 convolutional layers each with 24 kernel
- 128 unit LSTM for question processing, 32 unit word-lookup embedding
- 4-layer MLP consisting of 256 units per layer with ReLU non-linearities for g_θ .
- 3-layer MLP consisting of 256, 256 (with 50% dropout), and 29 units with ReLU non-linearities for f_ϕ .
- **No attention No ResNet**

- CNN- 4 convolutional layers each with 24 kernel
- 128 unit LSTM for question processing, 32 unit word-lookup embedding
- 4-layer MLP consisting of 256 units per layer with ReLU non-linearities for g_θ .
- 3-layer MLP consisting of 256, 256 (with 50% dropout), and 29 units with ReLU non-linearities for f_ϕ .
- **No attention No ResNet**
- Training was end-to-end with Adam optimizer with a learning rate of $2.5e^{-4}$

Relational networks

Model	Overall	Count	Exist	Compare numbers	Query	Compare attribute
<i>Human</i>	92.6	86.7	96.6	86.5	95.0	96.0
<i>LSTM</i>	46.8	41.7	61.1	69.8	36.8	51.8
<i>CNN+RNN</i>	52.3	43.7	65.2	67.1	49.3	53.0
<i>SA</i>	68.5	52.2	71.1	73.5	85.3	52.3
<i>NMN</i>	72.1	79.3	52.5	71.4	78.9	78
<i>N2NMN</i>	83.7	85.7	68.5	83.7	90	88.7
<i>PG</i>	96.9	97.1	92.7	98.6	98.1	98.9
<i>RN</i>	95.5	90.1	97.8	93.6	97.9	97.1

Performance on bAbI:

- Famous dataset for text-based QA.

Performance on bAbI:

- Famous dataset for text-based QA.
- There are 20 tasks, each corresponding to a particular type of reasoning.

Performance on bAbI:

- Famous dataset for text-based QA.
- There are 20 tasks, each corresponding to a particular type of reasoning.
- Each question is associated with a set of supporting facts.

Performance on bAbI:

- Famous dataset for text-based QA.
- There are 20 tasks, each corresponding to a particular type of reasoning.
- Each question is associated with a set of supporting facts.
- For example, the facts **Sandra picked up the football** and **Sandra went to the office** support the question **Where is the football?** (answer: **office**).

Performance on bAbI:

- Famous dataset for text-based QA.
- There are 20 tasks, each corresponding to a particular type of reasoning.
- Each question is associated with a set of supporting facts.
- For example, the facts **Sandra picked up the football** and **Sandra went to the office** support the question **Where is the football?** (answer: **office**).
- A model succeeds on a task if its performance surpasses 95%.

Performance on bAbI:

- Famous dataset for text-based QA.
- There are 20 tasks, each corresponding to a particular type of reasoning.
- Each question is associated with a set of supporting facts.
- For example, the facts **Sandra picked up the football** and **Sandra went to the office** support the question **Where is the football?** (answer: **office**).
- A model succeeds on a task if its performance surpasses 95%.
- Memory networks pass 14 out of 20.

Model:

- Identify up to 20 sentences in the support set that were immediately prior to the probe question.

Model:

- Identify up to 20 sentences in the support set that were immediately prior to the probe question.
- Tag these sentences with labels indicating their relative position in the support set, and processed each sentence word-by-word with an LSTM.

Model:

- Identify up to 20 sentences in the support set that were immediately prior to the probe question.
- Tag these sentences with labels indicating their relative position in the support set, and processed each sentence word-by-word with an LSTM.
- Previous bAbI models processed all word tokens from all support sentences sequentially.

Model:

- Identify up to 20 sentences in the support set that were immediately prior to the probe question.
- Tag these sentences with labels indicating their relative position in the support set, and processed each sentence word-by-word with an LSTM.
- Previous bAbI models processed all word tokens from all support sentences sequentially.
- The final state of the sentence-processing-LSTM is considered to be an object.

Model:

- Identify up to 20 sentences in the support set that were immediately prior to the probe question.
- Tag these sentences with labels indicating their relative position in the support set, and processed each sentence word-by-word with an LSTM.
- Previous bAbl models processed all word tokens from all support sentences sequentially.
- The final state of the sentence-processing-LSTM is considered to be an object.
- The model succeeded on 18 tasks.

Model:

- Identify up to 20 sentences in the support set that were immediately prior to the probe question.
- Tag these sentences with labels indicating their relative position in the support set, and processed each sentence word-by-word with an LSTM.
- Previous bAbl models processed all word tokens from all support sentences sequentially.
- The final state of the sentence-processing-LSTM is considered to be an object.
- The model succeeded on 18 tasks.
- Where other models failed miserably on these tasks, the model missed the 95% cut by 3.1% and 11.5%.

Performance on dynamic physical systems

- The paper introduces a dataset for 2 tasks on MuJoCo physics engine.

Performance on dynamic physical systems

- The paper introduces a dataset for 2 tasks on MuJoCo physics engine.
- Each scene contained 10 colored balls moving on a table-top surface.

Performance on dynamic physical systems

- The paper introduces a dataset for 2 tasks on MuJoCo physics engine.
- Each scene contained 10 colored balls moving on a table-top surface.
- Randomly selected ball pairs were connected by invisible springs or a rigid constraint. These connections prevented the balls from moving independently.

Performance on dynamic physical systems

- The paper introduces a dataset for 2 tasks on MuJoCo physics engine.
- Each scene contained 10 colored balls moving on a table-top surface.
- Randomly selected ball pairs were connected by invisible springs or a rigid constraint. These connections prevented the balls from moving independently.
- The paper introduces 2 tasks: We defined two separate tasks: 1) infer the existence or absence of connections between balls when only observing their color and coordinate positions across multiple sequential frames, and 2) count the number of systems on the table-top.

Performance on dynamic physical systems

- Both of these tasks involve reasoning about the relative positions and velocities of the balls to infer.

Performance on dynamic physical systems

- Both of these tasks involve reasoning about the relative positions and velocities of the balls to infer.
- In the connection inference task, our model correctly classified all the connections in 93% of the sample scenes in the test set.

Performance on dynamic physical systems

- Both of these tasks involve reasoning about the relative positions and velocities of the balls to infer.
- In the connection inference task, our model correctly classified all the connections in 93% of the sample scenes in the test set.
- In the counting task, the RN achieved similar performance, reporting the correct number of connected systems for 95% of the test scene samples.

Performance on dynamic physical systems

- Both of these tasks involve reasoning about the relative positions and velocities of the balls to infer.
- In the connection inference task, our model correctly classified all the connections in 93% of the sample scenes in the test set.
- In the counting task, the RN achieved similar performance, reporting the correct number of connected systems for 95% of the test scene samples.
- An MLP with comparable number of parameters was unable to perform better than chance for both tasks.

- This paper shows the difference between **processing** and **reasoning**.
Saying, ResNets are highly capable visual processors, but they may not be the most appropriate choice for reasoning about arbitrary relations.

- This paper shows the difference between **processing** and **reasoning**. Saying, ResNets are highly capable visual processors, but they may not be the most appropriate choice for reasoning about arbitrary relations.
- Demonstrate the RN's rich capacity for structured reasoning even with unstructured inputs and outputs.

- A model made from general-purpose components that could learn to visually reason, would likely be more widely applicable across domains.

- A model made from general-purpose components that could learn to visually reason, would likely be more widely applicable across domains.
- A FiLM layer carries out a simple, feature-wise affine transformation on a neural network's intermediate features, conditioned on an arbitrary input.

- A model made from general-purpose components that could learn to visually reason, would likely be more widely applicable across domains.
- A FiLM layer carries out a simple, feature-wise affine transformation on a neural network's intermediate features, conditioned on an arbitrary input.
- FiLM can be thought of as a generalization of Conditional Normalization, which has proven highly successful for image stylization.

- FiLM learns to adaptively influence the output of a neural network by applying an linear transformation.

- FiLM learns to adaptively influence the output of a neural network by applying an linear transformation.
- FiLM - Feature-wise Linearly Modulated

- FiLM learns to adaptively influence the output of a neural network by applying an linear transformation.
- FiLM - Feature-wise Linearly Modulated
- $\gamma_{i,c} = f_c(x_i)$

- FiLM learns to adaptively influence the output of a neural network by applying an linear transformation.
- FiLM - Feature-wise Linearly Modulated
- $\gamma_{i,c} = f_c(x_i)$
- $\beta_{i,c} = h_c(x_i)$

- FiLM learns to adaptively influence the output of a neural network by applying an linear transformation.
- FiLM - Feature-wise Linearly Modulated
- $\gamma_{i,c} = f_c(x_i)$
- $\beta_{i,c} = h_c(x_i)$
- $\gamma_{i,c}$ and $\beta_{i,c}$ modulate a neural network's activations $F_{i,c}$.

- FiLM learns to adaptively influence the output of a neural network by applying an linear transformation.
- FiLM - Feature-wise Linearly Modulated
- $\gamma_{i,c} = f_c(x_i)$
- $\beta_{i,c} = h_c(x_i)$
- $\gamma_{i,c}$ and $\beta_{i,c}$ modulate a neural network's activations $F_{i,c}$.
- $\text{FiLM}(F_{i,c}|\gamma_{i,c}, \beta_{i,c}) = \gamma_{i,c}F_{i,c} + \beta_{i,c}$

- FiLM learns to adaptively influence the output of a neural network by applying an linear transformation.
- FiLM - Feature-wise Linearly Modulated
- $\gamma_{i,c} = f_c(x_i)$
- $\beta_{i,c} = h_c(x_i)$
- $\gamma_{i,c}$ and $\beta_{i,c}$ modulate a neural network's activations $F_{i,c}$.
- $\text{FiLM}(F_{i,c}|\gamma_{i,c}, \beta_{i,c}) = \gamma_{i,c}F_{i,c} + \beta_{i,c}$
- FiLM layers empower the FiLM generator to manipulate feature maps of a target, FiLM-ed network by scaling them up or down, negating them, shutting them off, selectively thresholding them (when followed by a ReLU), and more.

- FiLM learns to adaptively influence the output of a neural network by applying an linear transformation.
- FiLM - Feature-wise Linearly Modulated
- $\gamma_{i,c} = f_c(x_i)$
- $\beta_{i,c} = h_c(x_i)$
- $\gamma_{i,c}$ and $\beta_{i,c}$ modulate a neural network's activations $F_{i,c}$.
- $\text{FiLM}(F_{i,c} | \gamma_{i,c}, \beta_{i,c}) = \gamma_{i,c} F_{i,c} + \beta_{i,c}$
- FiLM layers empower the FiLM generator to manipulate feature maps of a target, FiLM-ed network by scaling them up or down, negating them, shutting them off, selectively thresholding them (when followed by a ReLU), and more.
- Each feature map is conditioned independently, giving the FiLM generator moderately fine-grained control over activations at each FiLM layer.

- FiLM learns to adaptively influence the output of a neural network by applying an linear transformation.
- FiLM - Feature-wise Linearly Modulated
- $\gamma_{i,c} = f_c(x_i)$
- $\beta_{i,c} = h_c(x_i)$
- $\gamma_{i,c}$ and $\beta_{i,c}$ modulate a neural network's activations $F_{i,c}$.
- $\text{FiLM}(F_{i,c} | \gamma_{i,c}, \beta_{i,c}) = \gamma_{i,c} F_{i,c} + \beta_{i,c}$
- FiLM layers empower the FiLM generator to manipulate feature maps of a target, FiLM-ed network by scaling them up or down, negating them, shutting them off, selectively thresholding them (when followed by a ReLU), and more.
- Each feature map is conditioned independently, giving the FiLM generator moderately fine-grained control over activations at each FiLM layer.
- It is a scalable and computationally efficient conditioning method.

Model:

- Model consists of a FiLM-generating linguistic pipeline and a FiLM-ed visual pipeline.

Model:

- Model consists of a FiLM-generating linguistic pipeline and a FiLM-ed visual pipeline.
- The FiLM generator processes a question x_i using a Gated Recurrent Unit (GRU) network with 4096 hidden units that takes in learned, 200-dimensional word embeddings.

Model:

- Model consists of a FiLM-generating linguistic pipeline and a FiLM-ed visual pipeline.
- The FiLM generator processes a question x_i using a Gated Recurrent Unit (GRU) network with 4096 hidden units that takes in learned, 200-dimensional word embeddings.
- The final GRU hidden state is a question embedding, from which the model predicts (γ_i^n, β_i^n) for each n -th residual block.

Model:

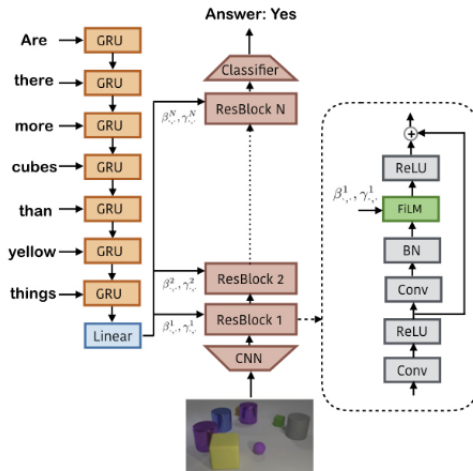
- Model consists of a FiLM-generating linguistic pipeline and a FiLM-ed visual pipeline.
- The FiLM generator processes a question x_i using a Gated Recurrent Unit (GRU) network with 4096 hidden units that takes in learned, 200-dimensional word embeddings.
- The final GRU hidden state is a question embedding, from which the model predicts (γ_i^n, β_i^n) for each n -th residual block.
- The visual pipeline extracts 128 14×14 image feature maps from a resized, 224×224 image input using CNN.

Model:

- Model consists of a FiLM-generating linguistic pipeline and a FiLM-ed visual pipeline.
- The FiLM generator processes a question x_i using a Gated Recurrent Unit (GRU) network with 4096 hidden units that takes in learned, 200-dimensional word embeddings.
- The final GRU hidden state is a question embedding, from which the model predicts (γ_i^n, β_i^n) for each n-th residual block.
- The visual pipeline extracts 128 14×14 image feature maps from a resized, 224×224 image input using CNN.
- The CNN architecture is similar to ResNet-101 with a added FiLM module. It was pretrained on Image-Net.


Model:

- Model consists of a FiLM-generating linguistic pipeline and a FiLM-ed visual pipeline.
- The FiLM generator processes a question x_i using a Gated Recurrent Unit (GRU) network with 4096 hidden units that takes in learned, 200-dimensional word embeddings.
- The final GRU hidden state is a question embedding, from which the model predicts (γ_i^n, β_i^n) for each n-th residual block.
- The visual pipeline extracts 128 14×14 image feature maps from a resized, 224×224 image input using CNN.
- The CNN architecture is similar to ResNet-101 with a added FiLM module. It was pretrained on Image-Net.
- Model is trained end-to-end from scratch with Adam (learning rate $3e-4$), weight decay ($1e-5$), batch size 64, and batch normalization and ReLU throughout FiLM-ed network.



- Images reveal that the FiLM model predicts using features areas near answer-related or question-related objects.


Q: *What shape is the...*



...purple thing? A: cube



...blue thing? A: sphere



Q: *How many cyan things are...*

...right of the gray cube?
A: 3

...left of the small cube?
A: 2

Model	Overall	Count	Exist	Compare numbers	Query	Compare attribute
<i>Human</i>	92.6	86.7	96.6	86.5	95.0	96.0
<i>LSTM</i>	46.8	41.7	61.1	69.8	36.8	51.8
<i>CNN+RNN</i>	52.3	43.7	65.2	67.1	49.3	53.0
<i>SA</i>	68.5	52.2	71.1	73.5	85.3	52.3
<i>NMN</i>	72.1	79.3	52.5	71.4	78.9	78
<i>N2NMN</i>	83.7	85.7	68.5	83.7	90	88.7
<i>PG</i>	96.9	97.1	92.7	98.6	98.1	98.9
<i>RN</i>	95.5	90.1	97.8	93.6	97.9	97.1
<i>FiLM</i>	97.7	94.3	99.1	96.8	99.1	99.1

We saw 3 different innovative ways to achieve super human performance in this task.

Thank You.